

#Predict the price of the Uber ride from a given pickup point to the agreed drop-off location. Perform following tasks:

1. Pre-process the dataset.
2. Identify outliers.
3. Check the correlation.
4. Implement linear regression and random forest regression models.
5. Evaluate the models and compare their respective scores like R2, RMSE, etc. Dataset link:
<https://www.kaggle.com/datasets/yasserh/uber-fares-dataset>
(<https://www.kaggle.com/datasets/yasserh/uber-fares-dataset>)

```
In [1]: #Importing the required libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: #importing the dataset
df = pd.read_csv("uber.csv")
```

1. Pre-process the dataset.

```
In [3]: df.head()
```

```
Out[3]:
```

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085

```
In [4]: df.info() #To get the required information of the dataset
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            200000 non-null  int64
1   key                   200000 non-null  object
2   fare_amount           200000 non-null  float64
3   pickup_datetime       200000 non-null  object
4   pickup_longitude      200000 non-null  float64
5   pickup_latitude       200000 non-null  float64
6   dropoff_longitude     199999 non-null  float64
7   dropoff_latitude      199999 non-null  float64
8   passenger_count       200000 non-null  int64
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

```
In [5]: df.columns #TO get number of columns in the dataset
```

```
Out[5]: Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
              'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
              'dropoff_latitude', 'passenger_count'],
              dtype='object')
```

```
In [6]: df = df.drop(['Unnamed: 0', 'key'], axis=1) #To drop unnamed column as it isn't
```

```
In [7]: df.head()
```

```
Out[7]:
```

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	40.7232
1	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	40.7500
2	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	40.7720
3	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	40.8030
4	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	40.7610

```
In [8]: df.shape #To get the total (Rows,Columns)
```

```
Out[8]: (200000, 7)
```

```
In [9]: df.dtypes #To get the type of each column
```

```
Out[9]: fare_amount      float64
pickup_datetime      object
pickup_longitude      float64
pickup_latitude      float64
dropoff_longitude      float64
dropoff_latitude      float64
passenger_count      int64
dtype: object
```

```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fare_amount            200000 non-null float64
1   pickup_datetime        200000 non-null object
2   pickup_longitude       200000 non-null float64
3   pickup_latitude        200000 non-null float64
4   dropoff_longitude      199999 non-null float64
5   dropoff_latitude       199999 non-null float64
6   passenger_count        200000 non-null int64
dtypes: float64(5), int64(1), object(1)
memory usage: 10.7+ MB
```

```
In [11]: df.describe() #To get statistics of each columns
```

```
Out[11]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
count	200000.000000	200000.000000	200000.000000	199999.000000	199999.000000	200000.000000
mean	11.359955	-72.527638	39.935885	-72.525292	39.923890	1.596901
std	9.901776	11.437787	7.720539	13.117408	6.794829	1.122449
min	-52.000000	-1340.648410	-74.015515	-3356.666300	-881.985513	0
25%	6.000000	-73.992065	40.734796	-73.991407	40.733823	1
50%	8.500000	-73.981823	40.752592	-73.980093	40.753042	1
75%	12.500000	-73.967154	40.767158	-73.963658	40.768001	1
max	499.000000	57.418457	1644.421482	1153.572603	872.697628	2

Filling Missing values

```
In [12]: df.isnull().sum()
```

```
Out[12]: fare_amount      0
pickup_datetime      0
pickup_longitude      0
pickup_latitude      0
dropoff_longitude      1
dropoff_latitude      1
passenger_count      0
dtype: int64
```

```
In [13]: df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(),inplace = True)
df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].median(),inplace = 1
```

```
In [14]: df.isnull().sum()
```

```
Out[14]: fare_amount      0
pickup_datetime      0
pickup_longitude      0
pickup_latitude      0
dropoff_longitude      0
dropoff_latitude      0
passenger_count      0
dtype: int64
```

```
In [15]: df.dtypes
```

```
Out[15]: fare_amount      float64
pickup_datetime      object
pickup_longitude      float64
pickup_latitude      float64
dropoff_longitude      float64
dropoff_latitude      float64
passenger_count      int64
dtype: object
```

Column pickup_datetime is in wrong format (Object). Convert it to DateTime Format

```
In [16]: df.pickup_datetime = pd.to_datetime(df.pickup_datetime, errors='coerce')
```

```
In [17]: df.dtypes
```

```
Out[17]: fare_amount      float64
pickup_datetime      datetime64[ns, UTC]
pickup_longitude      float64
pickup_latitude      float64
dropoff_longitude      float64
dropoff_latitude      float64
passenger_count      int64
dtype: object
```

To segregate each time of date and time

```
In [18]: df= df.assign(hour = df.pickup_datetime.dt.hour,  
                    day= df.pickup_datetime.dt.day,  
                    month = df.pickup_datetime.dt.month,  
                    year = df.pickup_datetime.dt.year,  
                    dayofweek = df.pickup_datetime.dt.dayofweek)
```

```
In [19]: df.head()
```

```
Out[19]:
```

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.999512	40.723217
1	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.994710	40.750325
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	40.772647
3	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	40.803349
4	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	40.761247

```
In [20]: # drop the column 'pickup_datetime' using drop()  
# 'axis = 1' drops the specified column  
  
df = df.drop('pickup_datetime',axis=1)
```

```
In [21]: df.head()
```

```
Out[21]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	7.5	-73.999817	40.738354	-73.999512	40.723217	1
1	7.7	-73.994355	40.728225	-73.994710	40.750325	1
2	12.9	-74.005043	40.740770	-73.962565	40.772647	1
3	5.3	-73.976124	40.790844	-73.965316	40.803349	1
4	16.0	-73.925023	40.744085	-73.973082	40.761247	1


```
In [22]: df.dtypes
```

```
Out[22]: fare_amount      float64  
pickup_longitude    float64  
pickup_latitude     float64  
dropoff_longitude   float64  
dropoff_latitude    float64  
passenger_count      int64  
hour                 int64  
day                  int64  
month                int64  
year                 int64  
dayofweek            int64  
dtype: object
```

Checking outliers and filling them

```
In [23]: df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20)) #Boxplot to
```

```
Out[23]: fare_amount      AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude    AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude      AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude    AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude      AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count      AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour                 AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day                 AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month               AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year               AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek           AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dtype: object
```



```
In [24]: #Using the InterQuartile Range to fill the values
def remove_outlier(df1 , col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_whisker = Q1-1.5*IQR
    upper_whisker = Q3+1.5*IQR
    df[col] = np.clip(df1[col] , lower_whisker , upper_whisker)
    return df1

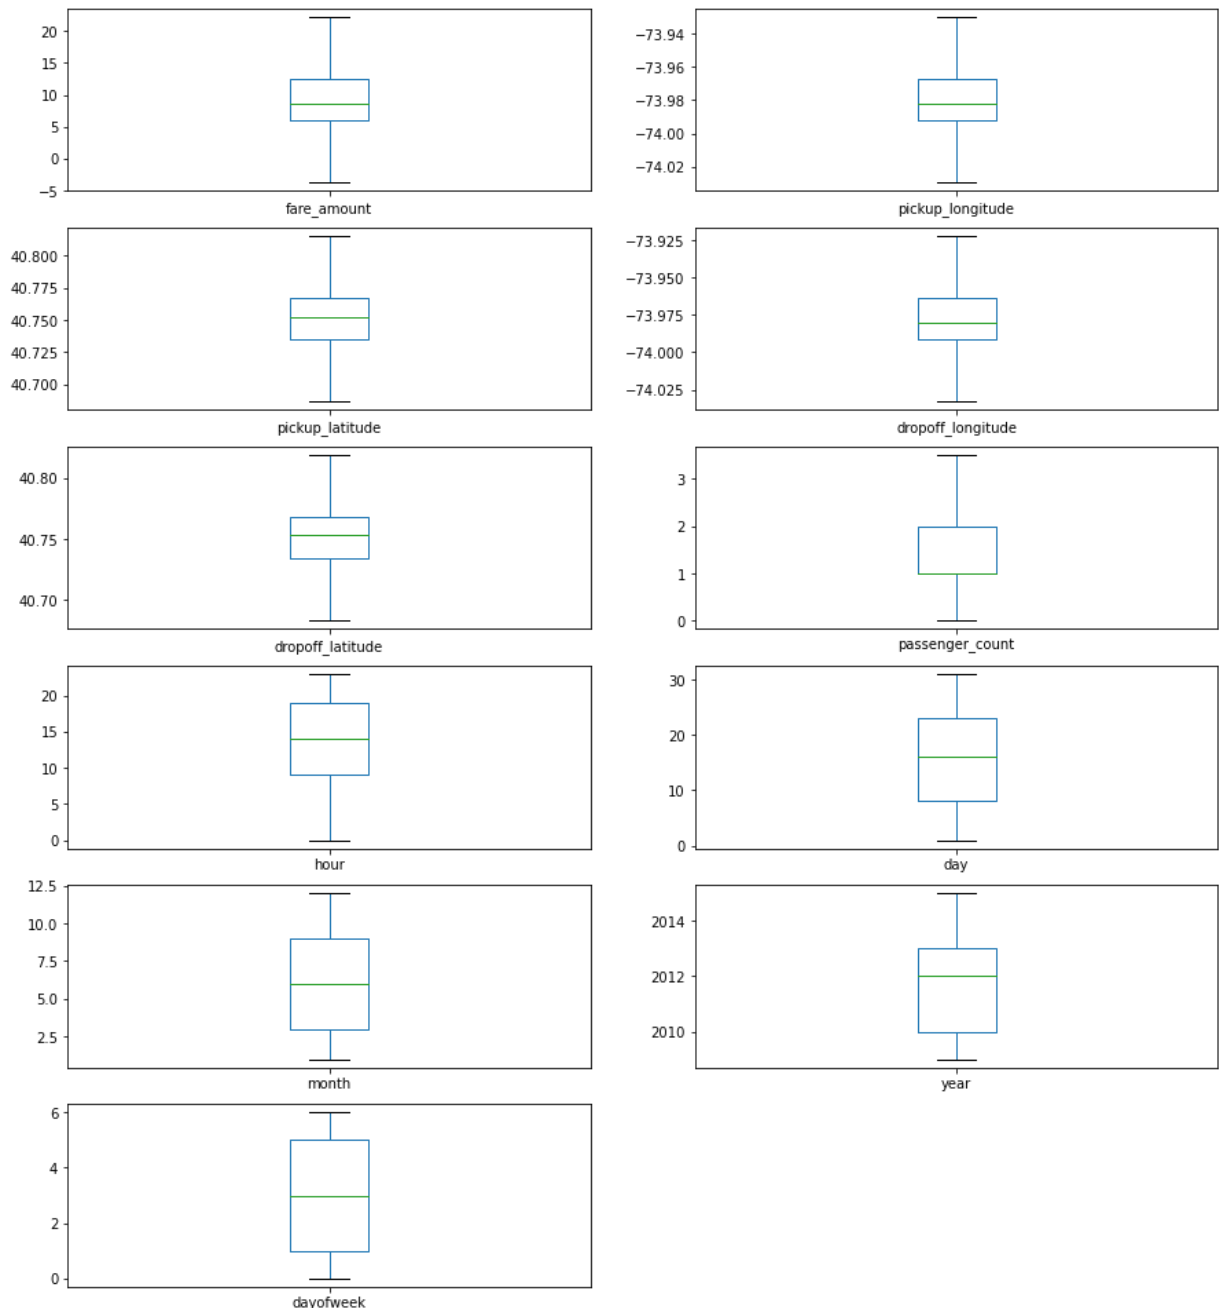
def treat_outliers_all(df1 , col_list):
    for c in col_list:
        df1 = remove_outlier(df , c)
    return df1
```

```
In [25]: df = treat_outliers_all(df , df.iloc[:, 0::])
```



```
In [26]: df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20)) #Boxplot sho
```

```
Out[26]: fare_amount      AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude    AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude      AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude    AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude      AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count      AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour                 AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day                 AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month               AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year               AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek           AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dtype: object
```



```
In [27]: #pip install haversine
import haversine as hs #Calculate the distance using Haversine to calculate the
travel_dist = []
for pos in range(len(df['pickup_longitude'])):
    long1,lati1,long2,lati2 = [df['pickup_longitude'][pos],df['pickup_latitude'][pos],df['dropoff_longitude'][pos],df['dropoff_latitude'][pos]]
    loc1=(lati1,long1)
    loc2=(lati2,long2)
    c = hs.haversine(loc1,loc2)
    travel_dist.append(c)

print(travel_dist)
df['dist_travel_km'] = travel_dist
df.head()
```

IOPub data rate exceeded.
The notebook server will temporarily stop sending output
to the client in order to avoid crashing it.
To change this limit, set the config variable
`--NotebookApp.iopub_data_rate_limit`.

Current values:
NotebookApp.iopub_data_rate_limit=1000000.0 (bytes/sec)
NotebookApp.rate_limit_window=3.0 (secs)

```
Out[27]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	7.5	-73.999817	40.738354	-73.999512	40.723217	1
1	7.7	-73.994355	40.728225	-73.994710	40.750325	1
2	12.9	-74.005043	40.740770	-73.962565	40.772647	1
3	5.3	-73.976124	40.790844	-73.965316	40.803349	1
4	16.0	-73.929786	40.744085	-73.973082	40.761247	1

```
In [28]: #Uber doesn't travel over 130 kms so minimize the distance
df= df.loc[(df.dist_travel_km >= 1) | (df.dist_travel_km <= 130)]
print("Remaining observastions in the dataset:", df.shape)
```

Remaining observastions in the dataset: (200000, 12)

```
In [29]: #Finding inccorect Latitude (Less than or greater than 90) and Longitude (greater
incorrect_coordinates = df.loc[(df.pickup_latitude > 90) | (df.pickup_latitude < -90) |
                                (df.dropoff_latitude > 90) | (df.dropoff_latitude < -90) |
                                (df.pickup_longitude > 180) | (df.pickup_longitude < -180) |
                                (df.dropoff_longitude > 180) | (df.dropoff_longitude < -180)]
```

```
In [30]: df.drop(incorrect_coordinates, inplace = True, errors = 'ignore')
```

```
In [31]: df.head()
```

```
Out[31]:
```


	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	7.5	-73.999817	40.738354	-73.999512	40.723217	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	
4	16.0	-73.929786	40.744085	-73.973082	40.761247	

```
In [32]: df.isnull().sum()
```

```
Out[32]: fare_amount      0
pickup_longitude    0
pickup_latitude     0
dropoff_longitude   0
dropoff_latitude    0
passenger_count     0
hour                0
day                 0
month               0
year                0
dayofweek           0
dist_travel_km      0
dtype: int64
```

```
In [33]: sns.heatmap(df.isnull()) #Free for null values
```

```
Out[33]: <AxesSubplot:>
```



```
In [34]: corr = df.corr() #Function to find the correlation
```

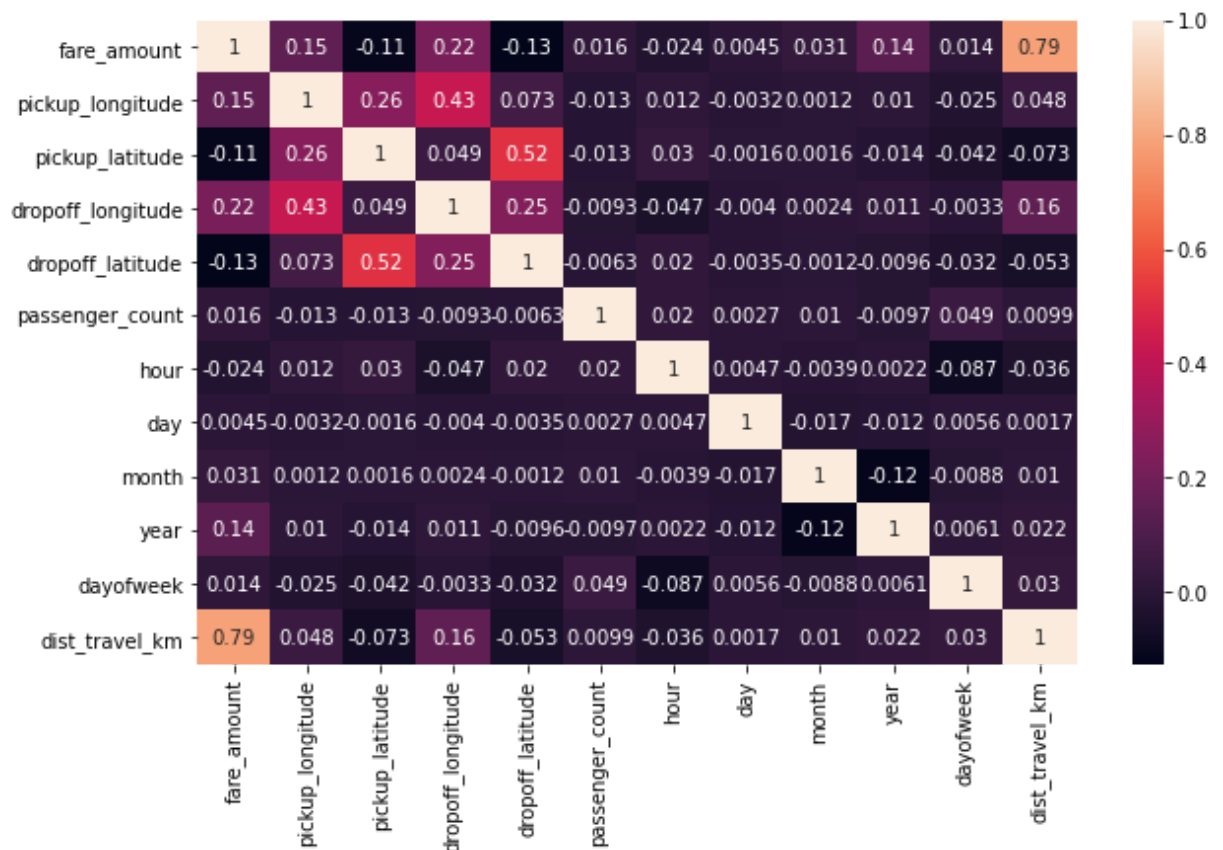
```
In [35]: corr
```

```
Out[35]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
fare_amount	1.000000	0.154069	-0.110842	0.218675	-0.12589
pickup_longitude	0.154069	1.000000	0.259497	0.425619	0.07329
pickup_latitude	-0.110842	0.259497	1.000000	0.048889	0.51571
dropoff_longitude	0.218675	0.425619	0.048889	1.000000	0.24566
dropoff_latitude	-0.125898	0.073290	0.515714	0.245667	1.00000
passenger_count	0.015778	-0.013213	-0.012889	-0.009303	-0.00630
hour	-0.023623	0.011579	0.029681	-0.046558	0.01978
day	0.004534	-0.003204	-0.001553	-0.004007	-0.00347
month	0.030817	0.001169	0.001562	0.002391	-0.00119
year	0.141277	0.010198	-0.014243	0.011346	-0.00960
dayofweek	0.013652	-0.024652	-0.042310	-0.003336	-0.03191
dist_travel_km	0.786385	0.048446	-0.073362	0.155191	-0.05270

```
In [36]: fig,axis = plt.subplots(figsize = (10,6))
sns.heatmap(df.corr(),annot = True) #Correlation Heatmap (Light values means high
```

```
Out[36]: <AxesSubplot:>
```



Dividing the dataset into feature and target values

```
In [182]: x = df[['pickup_longitude','pickup_latitude','dropoff_longitude','dropoff_latitude']]
```

```
In [183]: y = df['fare_amount']
```

Dividing the dataset into training and testing dataset

```
In [184]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(x,y,test_size = 0.33)
```

Linear Regression

```
In [185]: from sklearn.linear_model import LinearRegression
          regression = LinearRegression()
```

```
In [186]: regression.fit(X_train,y_train)
```

Out[186]: LinearRegression()

```
In [80]: regression.intercept_ #To find the linear intercept
```

Out[80]: 2640.1356169149753

```
In [187]: regression.coef_ #To find the linear coefficient
```

```
Out[187]: array([ 2.54805415e+01, -7.18365435e+00,  1.96232986e+01, -1.79401980e+01,
                  5.48472723e-02,  5.32910041e-03,  4.05930990e-03,  5.74261856e-02,
                  3.66574831e-01, -3.03753790e-02,  1.84233728e+00])
```

```
In [188]: prediction = regression.predict(X_test) #To predict the target values
```

```
In [189]: print(prediction)
```

```
[ 5.47848314 10.11016249 12.19490542 ... 7.11952609 20.2482979
 8.82791961]
```

```
In [190]: y_test
```

```
Out[190]: 155740      4.90
          47070      10.00
          116192     14.50
          164589      6.50
          154309     11.30
          ...
          76552      7.70
          27926      10.90
          38972      6.50
          120341     22.25
          178449      8.10
          Name: fare_amount, Length: 66000, dtype: float64
```

Metrics Evaluation using R2, Mean Squared Error, Root Mean Squared Error

```
In [191]: from sklearn.metrics import r2_score
```

```
In [192]: r2_score(y_test, prediction)
```

```
Out[192]: 0.6651880468683617
```

```
In [193]: from sklearn.metrics import mean_squared_error
```

```
In [194]: MSE = mean_squared_error(y_test, prediction)
```

```
In [195]: MSE
```

```
Out[195]: 9.961516917717704
```

```
In [196]: RMSE = np.sqrt(MSE)
```

```
In [197]: RMSE
```

```
Out[197]: 3.156187085348032
```

Random Forest Regression

```
In [198]: from sklearn.ensemble import RandomForestRegressor
```

```
In [199]: rf = RandomForestRegressor(n_estimators=100) #Here n_estimators means number of t
```

```
In [200]: rf.fit(X_train, y_train)
```

```
Out[200]: RandomForestRegressor()
```

```
In [201]: y_pred = rf.predict(X_test)
```

```
In [202]: y_pred
```

```
Out[202]: array([ 5.714 , 10.285 , 12.68  , ...,  6.338 , 19.4685,  7.712  ])
```

Metrics evaluation for Random Forest

```
In [210]: R2_Random = r2_score(y_test,y_pred)
```

```
In [211]: R2_Random
```

```
Out[211]: 0.7948374920410631
```

```
In [205]: MSE_Random = mean_squared_error(y_test,y_pred)
```

```
In [206]: MSE_Random
```

```
Out[206]: 6.104112397417331
```

```
In [207]: RMSE_Random = np.sqrt(MSE_Random)
```

```
In [208]: RMSE_Random
```

```
Out[208]: 2.4706501972997574
```