

Implement K-Means clustering/ hierarchical clustering on sales_data_sample.csv dataset. Determine the number of clusters using the elbow method. ¶

```
In [198]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
#Importing the required libraries.
```

```
In [199]: from sklearn.cluster import KMeans, k_means #For clustering
from sklearn.decomposition import PCA #Linear Dimensionality reduction.
```

```
In [200]: df = pd.read_csv("sales_data_sample.csv") #Loading the dataset.
```

Preprocessing

```
In [201]: df.head()
```

```
Out[201]:
```

| | ORDERNUMBER | QUANTITYORDERED | PRICEEACH | ORDERLINENUMBER | SALES | ORDERDATE |
|---|-------------|-----------------|-----------|-----------------|---------|-----------------|
| 0 | 10107 | 30 | 95.70 | 2 | 2871.00 | 2/24/2003 0:00 |
| 1 | 10121 | 34 | 81.35 | 5 | 2765.90 | 5/7/2003 0:00 |
| 2 | 10134 | 41 | 94.74 | 2 | 3884.34 | 7/1/2003 0:00 |
| 3 | 10145 | 45 | 83.26 | 6 | 3746.70 | 8/25/2003 0:00 |
| 4 | 10159 | 49 | 100.00 | 14 | 5205.27 | 10/10/2003 0:00 |

5 rows × 25 columns

```
In [202]: df.shape
```

```
Out[202]: (2823, 25)
```

```
In [203]: df.describe()
```

Out[203]:

| | ORDERNUMBER | QUANTITYORDERED | PRICEEACH | ORDERLINENUMBER | SALES | |
|-------|--------------|-----------------|-------------|-----------------|--------------|-----|
| count | 2823.000000 | 2823.000000 | 2823.000000 | 2823.000000 | 2823.000000 | 282 |
| mean | 10258.725115 | 35.092809 | 83.658544 | 6.466171 | 3553.889072 | |
| std | 92.085478 | 9.741443 | 20.174277 | 4.225841 | 1841.865106 | |
| min | 10100.000000 | 6.000000 | 26.880000 | 1.000000 | 482.130000 | |
| 25% | 10180.000000 | 27.000000 | 68.860000 | 3.000000 | 2203.430000 | |
| 50% | 10262.000000 | 35.000000 | 95.700000 | 6.000000 | 3184.800000 | |
| 75% | 10333.500000 | 43.000000 | 100.000000 | 9.000000 | 4508.000000 | |
| max | 10425.000000 | 97.000000 | 100.000000 | 18.000000 | 14082.800000 | |

```
In [204]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ORDERNUMBER           2823 non-null   int64
1   QUANTITYORDERED       2823 non-null   int64
2   PRICEEACH             2823 non-null   float64
3   ORDERLINENUMBER       2823 non-null   int64
4   SALES                 2823 non-null   float64
5   ORDERDATE             2823 non-null   object
6   STATUS               2823 non-null   object
7   QTR_ID               2823 non-null   int64
8   MONTH_ID             2823 non-null   int64
9   YEAR_ID              2823 non-null   int64
10  PRODUCTLINE           2823 non-null   object
11  MSRP                  2823 non-null   int64
12  PRODUCTCODE           2823 non-null   object
13  CUSTOMERNAME          2823 non-null   object
14  PHONE                 2823 non-null   object
15  ADDRESSLINE1          2823 non-null   object
16  ADDRESSLINE2          302 non-null    object
17  CITY                  2823 non-null   object
18  STATE                 1337 non-null   object
19  POSTALCODE            2747 non-null   object
20  COUNTRY               2823 non-null   object
21  TERRITORY             1749 non-null   object
22  CONTACTLASTNAME       2823 non-null   object
23  CONTACTFIRSTNAME      2823 non-null   object
24  DEALSIZE              2823 non-null   object
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB
```

```
In [205]: df.isnull().sum()
```

```
Out[205]: ORDERNUMBER          0
          QUANTITYORDERED      0
          PRICEEACH            0
          ORDERLINENUMBER      0
          SALES                 0
          ORDERDATE            0
          STATUS               0
          QTR_ID               0
          MONTH_ID             0
          YEAR_ID              0
          PRODUCTLINE          0
          MSRP                 0
          PRODUCTCODE          0
          CUSTOMERNAME         0
          PHONE                0
          ADDRESSLINE1         0
          ADDRESSLINE2        2521
          CITY                 0
          STATE                1486
          POSTALCODE           76
          COUNTRY              0
          TERRITORY            1074
          CONTACTLASTNAME      0
          CONTACTFIRSTNAME     0
          DEALSIZE             0
          dtype: int64
```

```
In [206]: df.dtypes
```

```
Out[206]: ORDERNUMBER          int64
QUANTITYORDERED        int64
PRICEEACH              float64
ORDERLINENUMBER        int64
SALES                  float64
ORDERDATE              object
STATUS                 object
QTR_ID                 int64
MONTH_ID               int64
YEAR_ID                int64
PRODUCTLINE            object
MSRP                   int64
PRODUCTCODE            object
CUSTOMERNAME           object
PHONE                  object
ADDRESSLINE1           object
ADDRESSLINE2           object
CITY                   object
STATE                  object
POSTALCODE             object
COUNTRY                object
TERRITORY              object
CONTACTLASTNAME        object
CONTACTFIRSTNAME       object
DEALSIZE               object
dtype: object
```

```
In [207]: df_drop = ['ADDRESSLINE1', 'ADDRESSLINE2', 'STATUS', 'POSTALCODE', 'CITY', 'TERRITORY']
df = df.drop(df_drop, axis=1) #Dropping the categorical unnecessary columns along axis=1
```

```
In [208]: df.isnull().sum()
```

```
Out[208]: QUANTITYORDERED    0
PRICEEACH                    0
ORDERLINENUMBER              0
SALES                        0
ORDERDATE                    0
QTR_ID                       0
MONTH_ID                     0
YEAR_ID                      0
PRODUCTLINE                  0
MSRP                         0
PRODUCTCODE                  0
COUNTRY                      0
DEALSIZE                     0
dtype: int64
```

```
In [209]: df.dtypes
```

```
Out[209]: QUANTITYORDERED    int64
PRICEEACH                    float64
ORDERLINENUMBER              int64
SALES                        float64
ORDERDATE                    object
QTR_ID                       int64
MONTH_ID                     int64
YEAR_ID                      int64
PRODUCTLINE                  object
MSRP                         int64
PRODUCTCODE                  object
COUNTRY                       object
DEALSIZE                      object
dtype: object
```

```
In [ ]: # Checking the categorical columns.
```

```
In [210]: df['COUNTRY'].unique()
```

```
Out[210]: array(['USA', 'France', 'Norway', 'Australia', 'Finland', 'Austria', 'UK',
                'Spain', 'Sweden', 'Singapore', 'Canada', 'Japan', 'Italy',
                'Denmark', 'Belgium', 'Philippines', 'Germany', 'Switzerland',
                'Ireland'], dtype=object)
```

```
In [211]: df['PRODUCTLINE'].unique()
```

```
Out[211]: array(['Motorcycles', 'Classic Cars', 'Trucks and Buses', 'Vintage Cars',
                'Planes', 'Ships', 'Trains'], dtype=object)
```

```
In [212]: df['DEALSIZE'].unique()
```

```
Out[212]: array(['Small', 'Medium', 'Large'], dtype=object)
```

```
In [213]: productline = pd.get_dummies(df['PRODUCTLINE']) #Converting the categorical column
Dealsize = pd.get_dummies(df['DEALSIZE'])
```

```
In [214]: df = pd.concat([df,productline,Dealsize], axis = 1)
```

```
In [215]: df_drop = ['COUNTRY','PRODUCTLINE','DEALSIZE'] #Dropping Country too as there are no more
df = df.drop(df_drop, axis=1)
```

```
In [216]: df['PRODUCTCODE'] = pd.Categorical(df['PRODUCTCODE']).codes #Converting the data to
```

```
In [217]: df.drop('ORDERDATE', axis=1, inplace=True) #Dropping the Orderdate as Month is al
```


```
In [218]: df.dtypes #All the datatypes are converted into numeric
```

```
Out[218]: QUANTITYORDERED      int64
PRICEEACH      float64
ORDERLINENUMBER  int64
SALES      float64
QTR_ID      int64
MONTH_ID      int64
YEAR_ID      int64
MSRP      int64
PRODUCTCODE      int8
Classic Cars      uint8
Motorcycles      uint8
Planes      uint8
Ships      uint8
Trains      uint8
Trucks and Buses      uint8
Vintage Cars      uint8
Large      uint8
Medium      uint8
Small      uint8
dtype: object
```

Plotting the Elbow Plot to determine the number of clusters.

```
In [219]: distortions = [] # Within Cluster Sum of Squares from the centroid
K = range(1,10)
for k in K:
    kmeanModel = KMeans(n_clusters=k)
    kmeanModel.fit(df)
    distortions.append(kmeanModel.inertia_) #Appending the inertia to the Disto
```

```
In [220]: plt.figure(figsize=(16,8))
plt.plot(K, distortions, 'bx-')
plt.xlabel('k')
plt.ylabel('Distortion')
plt.title('The Elbow Method showing the optimal k')
plt.show()
```



As the number of k increases Inertia decreases.

Observations: A Elbow can be observed at 3 and after that the curve decreases gradually.

```
In [221]: X_train = df.values #Returns a numpy array.
```

```
In [222]: X_train.shape
```

```
Out[222]: (2823, 19)
```

```
In [223]: model = KMeans(n_clusters=3,random_state=2) #Number of cluster = 3
model = model.fit(X_train) #Fitting the values to create a model.
predictions = model.predict(X_train) #Predicting the cluster values (0,1,or 2)
```

```
In [225]: unique,counts = np.unique(predictions,return_counts=True)
```

```
In [226]: counts = counts.reshape(1,3)
```

```
In [227]: counts_df = pd.DataFrame(counts,columns=['Cluster1','Cluster2','Cluster3'])
```

```
In [228]: counts_df.head()
```

```
Out[228]:
```

| | Cluster1 | Cluster2 | Cluster3 |
|---|----------|----------|----------|
| 0 | 1083 | 1367 | 373 |

Visualization

```
In [229]: pca = PCA(n_components=2) #Converting all the features into 2 columns to make it
```

```
In [230]: reduced_X = pd.DataFrame(pca.fit_transform(X_train),columns=['PCA1','PCA2']) #Cre
```


```
In [231]: reduced_X.head()
```

```
Out[231]:
```

| | PCA1 | PCA2 |
|---|-------------|------------|
| 0 | -682.488323 | -42.819535 |
| 1 | -787.665502 | -41.694991 |
| 2 | 330.732170 | -26.481208 |
| 3 | 193.040232 | -26.285766 |
| 4 | 1651.532874 | -6.891196 |


```
In [232]: #Plotting the normal Scatter Plot
plt.figure(figsize=(14,10))
plt.scatter(reduced_X['PCA1'],reduced_X['PCA2'])
```

```
Out[232]: <matplotlib.collections.PathCollection at 0x218dc747880>
```



```
In [233]: model.cluster_centers_ #Finding the centriods. (3 Centriods in total. Each Array
```

```
Out[233]: array([[ 3.72031394e+01,  9.52120960e+01,  6.44967682e+00,
  4.13868425e+03,  2.72022161e+00,  7.09879963e+00,
  2.00379409e+03,  1.13248384e+02,  5.04469067e+01,
  3.74884580e-01,  1.15420129e-01,  9.41828255e-02,
  8.21791320e-02,  1.84672207e-02,  1.16343490e-01,
  1.98522622e-01,  2.08166817e-17,  1.00000000e+00,
 -6.66133815e-16],
 [ 3.08302853e+01,  7.00755230e+01,  6.67300658e+00,
  2.12409474e+03,  2.71762985e+00,  7.09509876e+00,
  2.00381127e+03,  7.84784199e+01,  6.24871982e+01,
  2.64813460e-01,  1.21433797e-01,  1.29480614e-01,
  1.00219459e-01,  3.87710315e-02,  9.21726408e-02,
  2.53108998e-01,  6.93889390e-18,  6.21799561e-02,
  9.37820044e-01],
 [ 4.45871314e+01,  9.98931099e+01,  5.75603217e+00,
  7.09596863e+03,  2.71045576e+00,  7.06434316e+00,
  2.00389008e+03,  1.45823056e+02,  3.14959786e+01,
  5.33512064e-01,  1.07238606e-01,  7.23860590e-02,
  2.14477212e-02,  1.07238606e-02,  1.31367292e-01,
  1.23324397e-01,  4.20911528e-01,  5.79088472e-01,
  5.55111512e-17]])
```


```
In [234]: reduced_centers = pca.transform(model.cluster_centers_) #Transforming the centroids
```

```
In [235]: reduced_centers
```

```
Out[235]: array([[ 5.84994044e+02, -4.36786931e+00],  
                [-1.43005891e+03,  2.60041009e+00],  
                [ 3.54247180e+03,  3.15185487e+00]])
```

```
In [236]: plt.figure(figsize=(14,10))  
plt.scatter(reduced_X['PCA1'],reduced_X['PCA2'])  
plt.scatter(reduced_centers[:,0],reduced_centers[:,1],color='black',marker='x',s=
```

```
Out[236]: <matplotlib.collections.PathCollection at 0x218deb6e220>
```



```
In [237]: reduced_X['Clusters'] = predictions #Adding the Clusters to the reduced dataframe
```

```
In [238]: reduced_X.head()
```

```
Out[238]:
```

| | PCA1 | PCA2 | Clusters |
|---|-------------|------------|----------|
| 0 | -682.488323 | -42.819535 | 1 |
| 1 | -787.665502 | -41.694991 | 1 |
| 2 | 330.732170 | -26.481208 | 0 |
| 3 | 193.040232 | -26.285766 | 0 |
| 4 | 1651.532874 | -6.891196 | 0 |

```

In [239]: #Plotting the clusters
plt.figure(figsize=(14,10))
#               taking the cluster number and first column               taking
plt.scatter(reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA2'], color='red', marker='o')
plt.scatter(reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA2'], color='blue', marker='o')
plt.scatter(reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA2'], color='green', marker='o')

plt.scatter(reduced_centers[:,0], reduced_centers[:,1], color='black', marker='x', s=100)

```

Out[239]: <matplotlib.collections.PathCollection at 0x218dce9e1f0>

